

ecommerce_data_peerreview

December 14, 2023

0.1 Importing Libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import r2_score
from catboost import CatBoostRegressor
```

0.2 Converting Data to Pandas DataFrame

```
[2]: df = pd.read_csv('Ecommerce Customers.csv')
```

0.3 Reading the first 5 rows

```
[3]: df.head()
```

```
[3]:
```

	Email	\
0	mstephenson@fernandez.com	
1	hduke@hotmail.com	
2	pallen@yahoo.com	
3	riverarebecca@gmail.com	
4	mstephens@davidson-herman.com	

	Address	Avatar	\
0	835 Frank Tunnel\nWrightmouth, MI 82180-9605	Violet	
1	4547 Archer Common\nDiazchester, CA 06566-8576	DarkGreen	
2	24645 Valerie Unions Suite 582\nCobbborough, D...	Bisque	
3	1414 David Throughway\nPort Jason, OH 22070-1220	SaddleBrown	
4	14023 Rodriguez Passage\nPort Jacobville, PR 3...	MediumAquaMarine	

	Avg. Session Length	Time on App	Time on Website	Length of Membership	\
--	---------------------	-------------	-----------------	----------------------	---

0	34.497268	12.655651	39.577668	4.082621
1	31.926272	11.109461	37.268959	2.664034
2	33.000915	11.330278	37.110597	4.104543
3	34.305557	13.717514	36.721283	3.120179
4	33.330673	12.795189	37.536653	4.446308

	Yearly Amount Spent
0	587.951054
1	392.204933
2	487.547505
3	581.852344
4	599.406092

0.4 Getting data information using describe and info

```
[4]: df.describe()
```

```
[4]:
```

	Avg. Session Length	Time on App	Time on Website \
count	500.000000	500.000000	500.000000
mean	33.053194	12.052488	37.060445
std	0.992563	0.994216	1.010489
min	29.532429	8.508152	33.913847
25%	32.341822	11.388153	36.349257
50%	33.082008	11.983231	37.069367
75%	33.711985	12.753850	37.716432
max	36.139662	15.126994	40.005182

	Length of Membership	Yearly Amount Spent
count	500.000000	500.000000
mean	3.533462	499.314038
std	0.999278	79.314782
min	0.269901	256.670582
25%	2.930450	445.038277
50%	3.533975	498.887875
75%	4.126502	549.313828
max	6.922689	765.518462


```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 8 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Email               500 non-null    object
1   Address             500 non-null    object
2   Avatar              500 non-null    object
3   Avg. Session Length 500 non-null    float64
```


```
4   Time on App          500 non-null   float64
5   Time on Website      500 non-null   float64
6   Length of Membership  500 non-null   float64
7   Yearly Amount Spent  500 non-null   float64
dtypes: float64(5), object(3)
memory usage: 31.4+ KB
```

0.5 Plotting the data


```
[6]: sns.jointplot(data=df, x="Time on Website", y="Yearly Amount Spent", kind="reg")
plt.show()
```




```
[7]: sns.jointplot(data=df, x="Time on App", y="Length of Membership", kind="reg")  
plt.show()
```



```
[8]: sns.pairplot(data=df, diag_kind = 'kde')  
plt.show()
```



```
[9]: sns.regplot(data=df, x="Length of Membership", y="Yearly Amount Spent", ci=None)
plt.show()
```



0.6 Training the dataset

```
[10]: X = df.drop(['Yearly Amount Spent', 'Email', 'Address', 'Avatar'], axis=1)
      y = df['Yearly Amount Spent']
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

0.7 Finding prediction, coefficients and score using LinearRegression()

```
[11]: lr = LinearRegression().fit(X_train,y_train)
```

```
[12]: pd.DataFrame(zip(X.columns, lr.coef_))
```

```
[12]:
```

		0	1
0	Avg. Session Length	25.540333	
1	Time on App	38.756677	
2	Time on Website	0.095358	
3	Length of Membership	61.395478	

0.7.1 Clearly the time on app has a much higher coefficient than the time on the website, meaning it influences a lot more our target. A focus on the app can help increase the values.

```
[13]: lr.intercept_
```

```
[13]: -1032.4048272666016
```

```
[14]: y_pred = lr.predict(X_test)
      lr.score(X_test, y_test)
```

```
[14]: 0.9862104585031302
```

```
[15]: print('MSE: ', mean_squared_error(y_test,y_pred))
      print('MAE: ', mean_absolute_error(y_test,y_pred))
      print('RMSE: ', (np.sqrt(mean_squared_error(y_test,y_pred))))
      print('R2 Score:' , str(round(r2_score(y_test, y_pred) * 100)) + '%')
```

```
MSE: 84.17911922645058
MAE: 7.450059825799415
RMSE: 9.174917941128987
R2 Score: 99%
```

0.8 Finding scores using CatBoostRegressor

```
[21]: model = CatBoostRegressor(iterations = 6542, learning_rate = 0.02,
      ↪loss_function = 'RMSE')
```

```
[33]: # model.fit(X_train,y_train,eval_set= (X_test,y_test),) hiding results (too
      ↪long)
```

```
[29]: pred = model.predict(X_test)
      rmse = (np.sqrt(mean_squared_error(y_test,pred)))
      r2 = r2_score(y_test,pred)
```


```
[30]: print('MSE: ', mean_squared_error(y_test,pred))
      print('MAE: ', mean_absolute_error(y_test,pred))
      print('RMSE: ', rmse)
      print('R2 Score:' , str(round(r2_score(y_test, pred) * 100)) + '%')
```

```
MSE: 219.41178083745407
MAE: 10.725174838054881
RMSE: 14.812554838293565
R2 Score: 96%
```

0.9 Plotting the scores

```
[31]: plt.hist(y_test, bins=30,color= 'b')  
      plt.hist(pred, bins=30, color= 'r')
```

```
[31]: (array([2., 3., 4., 3., 6., 5., 6., 3., 9., 1., 6., 7., 5., 4., 7., 6., 5.,  
             1., 2., 2., 2., 1., 2., 2., 0., 2., 1., 0., 2., 1.]),  
      array([384.02256739, 394.44480081, 404.86703422, 415.28926764,  
            425.71150105, 436.13373447, 446.55596788, 456.9782013 ,  
            467.40043471, 477.82266813, 488.24490154, 498.66713496,  
            509.08936837, 519.51160179, 529.9338352 , 540.35606862,  
            550.77830203, 561.20053545, 571.62276886, 582.04500228,  
            592.46723569, 602.8894691 , 613.31170252, 623.73393593,  
            634.15616935, 644.57840276, 655.00063618, 665.42286959,  
            675.84510301, 686.26733642, 696.68956984]),  
      <BarContainer object of 30 artists>)
```



```
[32]: plt.hist(y_test, bins=30, color='b')  
      plt.hist(y_pred, bins=30, color= 'r')
```


```
[32]: (array([ 2.,  0.,  4.,  5.,  7.,  4.,  7.,  8.,  6.,  5., 10.,  5.,  5.,  
             8.,  6.,  3.,  2.,  2.,  3.,  1.,  2.,  0.,  1.,  1.,  1.,  0.,
```



```

1., 0., 0., 1.]),
array([366.65965982, 379.95435107, 393.24904231, 406.54373356,
419.83842481, 433.13311606, 446.42780731, 459.72249856,
473.0171898 , 486.31188105, 499.6065723 , 512.90126355,
526.1959548 , 539.49064605, 552.78533729, 566.08002854,
579.37471979, 592.66941104, 605.96410229, 619.25879354,
632.55348478, 645.84817603, 659.14286728, 672.43755853,
685.73224978, 699.02694102, 712.32163227, 725.61632352,
738.91101477, 752.20570602, 765.50039727]),
<BarContainer object of 30 artists>)

```



[]: