

Supervised Learning Practice


September 21, 2023

```
[32]: %matplotlib inline
import numpy as np
import pandas as pd
from pylab import *
import matplotlib.pyplot as plt
```

```
[8]: np.random.seed(2)
```

```
[9]: pageSpeeds= np.random.normal(3.0, 1.0, 100)
purchaseAmount= np.random.normal(50.0, 30.0, 100)/pageSpeeds
scatter(pageSpeeds, purchaseAmount)
```

```
[9]: <matplotlib.collections.PathCollection at 0x285eb647e80>
```




```
[12]: trainX=pageSpeeds[:80]
      testX= pageSpeeds[80:]

      trainY= purchaseAmount[:80]
      testY = purchaseAmount[80:]


      scatter(trainX,trainY)
      scatter(testX,testY)
```

[12]: <matplotlib.collections.PathCollection at 0x285eb776850>



```
[13]: scatter(testX,testY)
```


[13]: <matplotlib.collections.PathCollection at 0x285eb8f3be0>




```
[15]: x = np.array(trainX)
      y = np.array(trainY)

      p4 = np.poly1d(np.polyfit(x, y, 8))
```


```
[17]: xp = linspace(0, 7, 100)
      axes = plt.axes()
      axes.set_xlim([0,7])
      axes.set_ylim([0, 200])
      plt.scatter(x,y)
      plt.plot(xp, p4(xp), c='r')
      plt.show()
```




```
[31]: p4 = np.poly1d(np.polyfit(x, y, 5))
xp = linspace(0, 7, 100)
axes = plt.axes()
axes.set_xlim([0,7])
axes.set_ylim([0, 200])
plt.scatter(x,y)
plt.plot(xp, p4(xp), c='r')
plt.show()
```




```
[33]: p4 = np.poly1d(np.polyfit(x, y, 7))
xp = linspace(0, 7, 100)
axes = plt.axes()
axes.set_xlim([0,7])
axes.set_ylim([0, 200])
plt.scatter(x,y)
plt.plot(xp, p4(xp), c='r')
plt.show()
```




```
[20]: p4 = np.poly1d(np.polyfit(x, y, 6))
xp = linspace(0, 7, 100)
axes = plt.axes()
axes.set_xlim([0,7])
axes.set_ylim([0, 200])
plt.scatter(x,y)
plt.plot(xp, p4(xp), c='r')
plt.show()
```




```
[21]: p4 = np.poly1d(np.polyfit(x, y, 1))
xp = linspace(0, 7, 100)
axes = plt.axes()
axes.set_xlim([0,7])
axes.set_ylim([0, 200])
plt.scatter(x,y)
plt.plot(xp, p4(xp), c='r')
plt.show()
```



```
[29]: p4 = np.poly1d(np.polyfit(x, y, 2))
xp = linspace(0, 7, 100)
axes = plt.axes()
axes.set_xlim([0,7])
axes.set_ylim([0, 200])
plt.scatter(x,y)
plt.plot(xp, p4(xp), c='r')
plt.show()
```

```
[23]: p4 = np.poly1d(np.polyfit(x, y, 4))
xp = linspace(0, 7, 100)
axes = plt.axes()
axes.set_xlim([0,7])
axes.set_ylim([0, 200])
plt.scatter(x,y)
plt.plot(xp, p4(xp), c='r')
plt.show()
```



```
[8]: from sklearn.metrics import r2_score
r2 = r2_score(testY, p4(testX))
print(r2)

### best fit with 7th grade polynomial
```

```
-----
NameError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_5976\1044721873.py in <module>
      1 from sklearn.metrics import r2_score
----> 2 r2 = r2_score(testY, p4(testX))
      3 print(r2)
      4
      5 ### best fit with 7th grade polynomial


NameError: name 'testY' is not defined
```

```
[35]: r2 = r2_score(np.array(trainY), p4(np.array(trainX)))
print(r2)
```


0.6170116571732289

```
[3]: %matplotlib inline
from pylab import *
np.random.seed(2)
pageSpeeds = np.random.normal(3,1,1000)
purchaseAmount = np.random.normal(50,10,1000)/pageSpeeds
scatter(pageSpeeds, purchaseAmount)
```

[3]: <matplotlib.collections.PathCollection at 0x2787ae41f70>



```
[29]: x = np.array(pageSpeeds)
y = np.array(purchaseAmount)
p4 = np.poly1d(np.polyfit(x,y,5))
xp = np.linspace(0, 7, 100)
plt.scatter(x,y)
plt.plot(xp,p4(xp), c='r')
plt.show()
```



```
[30]: r2 = r2_score(y, p4(x))
      print(r2)
```

0.8553884386186104

```
[34]: df = pd.read_excel('http://cdn.sundog-soft.com/Udemy/DataScience/cars.xls')
      df.head()
```

```
[34]:
```

	Price	Mileage	Make	Model	Trim	Type	Cylinder	Liter	\
0	17314.103129	8221	Buick	Century	Sedan 4D	Sedan	6	3.1	
1	17542.036083	9135	Buick	Century	Sedan 4D	Sedan	6	3.1	
2	16218.847862	13196	Buick	Century	Sedan 4D	Sedan	6	3.1	
3	16336.913140	16342	Buick	Century	Sedan 4D	Sedan	6	3.1	
4	16339.170324	19832	Buick	Century	Sedan 4D	Sedan	6	3.1	

	Doors	Cruise	Sound	Leather
0	4	1	1	1
1	4	1	1	0
2	4	1	1	0
3	4	1	0	0
4	4	1	0	1

```
[54]: df
```

```
[54]:
```


	Price	Mileage	Make	Model	Trim	Type	Cylinder	\
0	17314.103129	8221	Buick	Century	Sedan 4D	Sedan	6	
1	17542.036083	9135	Buick	Century	Sedan 4D	Sedan	6	
2	16218.847862	13196	Buick	Century	Sedan 4D	Sedan	6	
3	16336.913140	16342	Buick	Century	Sedan 4D	Sedan	6	
4	16339.170324	19832	Buick	Century	Sedan 4D	Sedan	6	
..	
799	16507.070267	16229	Saturn	L Series	L300 Sedan 4D	Sedan	6	
800	16175.957604	19095	Saturn	L Series	L300 Sedan 4D	Sedan	6	
801	15731.132897	20484	Saturn	L Series	L300 Sedan 4D	Sedan	6	
802	15118.893228	25979	Saturn	L Series	L300 Sedan 4D	Sedan	6	
803	13585.636802	35662	Saturn	L Series	L300 Sedan 4D	Sedan	6	

	Liter	Doors	Cruise	Sound	Leather	Model_ord
0	3.1	4	1	1	1	10
1	3.1	4	1	1	0	10
2	3.1	4	1	1	0	10
3	3.1	4	1	0	0	10
4	3.1	4	1	0	1	10
..
799	3.0	4	1	0	0	21
800	3.0	4	1	1	0	21
801	3.0	4	1	1	0	21
802	3.0	4	1	1	0	21
803	3.0	4	1	0	0	21

```
[804 rows x 13 columns]
```

```
[149]: import seaborn as sns
import matplotlib.pyplot as plt
sns.scatterplot(x = df.Mileage, y = df.Price, hue = df.Doors, palette =_
↳ ['darkgreen', 'orange'])
```

```
[149]: <AxesSubplot:xlabel='Mileage', ylabel='Price'>
```



```
[53]: df.isnull().sum()
```

```
[53]: Price      0
      Mileage    0
      Make      0
      Model     0
      Trim      0
      Type      0
      Cylinder   0
      Liter      0
      Doors     0
      Cruise    0
      Sound     0
      Leather   0
      Model_ord  0
      dtype: int64
```

```
[74]: cylinders_df = df.groupby('Cylinder').count()
      cylinders_df
```

```
[74]:
```

	Price	Mileage	Make	Model	Trim	Type	Liter	Doors	Cruise	\
Cylinder										
6	804	804	804	804	804	804	804	804	804	

	Sound	Leather	Model_ord
Cylinder			
6	804	804	804

```
[80]: make_df = df.groupby('Make').mean()
make_df
```

```
[80]:
```

	Price	Mileage	Cylinder	Liter	Doors	Cruise	\
Make							
Buick	20815.113883	20428.100000	6.0	3.662500	4.000000	1.000000	
Cadillac	40936.335448	18908.562500	6.0	4.387500	3.750000	1.000000	
Chevrolet	16427.599348	19655.587500	6.0	2.868750	3.375000	0.596875	
Pontiac	18412.100422	19320.660000	6.0	3.300000	3.600000	0.753333	
SAAB	29494.704687	20964.122807	6.0	2.149123	3.473684	1.000000	
Saturn	13978.807560	20335.750000	6.0	2.333333	3.333333	0.450000	

	Sound	Leather	Model_ord
Make			
Buick	0.662500	0.437500	21.750000
Cadillac	0.550000	1.000000	17.875000
Chevrolet	0.828125	0.809375	14.281250
Pontiac	0.606667	0.640000	18.066667
SAAB	0.578947	0.728070	2.368421
Saturn	0.450000	0.483333	20.166667

```
[141]: df['Make'].value_counts()
```

```
[141]: Chevrolet    320
Pontiac          150
SAAB             114
Buick             80
Cadillac          80
Saturn            60
Name: Make, dtype: int64
```

```
[142]: df['Model'].value_counts()
```

```
[142]: Malibu        60
AVEO              60
Cavalier          60
Ion               50
Cobalt            50
9_3 H0           40
Vibe              30
```

```

Bonneville      30
Monte Carlo     30
Lacrosse        30
Impala          30
Grand Prix      30
9_5             30
Deville         30
Lesabre         20
Corvette        20
9_3             20
9_5 HO          20
G6              20
Grand Am        20
Park Avenue     20
Sunfire         10
Century         10
GTO             10
Classic         10
XLR-V8          10
STS-V8          10
STS-V6          10
CTS             10
CST-V           10
L Series        10
9-2X AWD        4
Name: Model, dtype: int64

```

```
df.describe().round(2)
```

```
[39]: import statsmodels.api as sm
```

```
[42]: df['Model_ord'] = pd.Categorical(df.Model).codes
df.head()
```

```
[42]:
```

	Price	Mileage	Make	Model	Trim	Type	Cylinder	Liter	\
0	17314.103129	8221	Buick	Century	Sedan 4D	Sedan	6	3.1	
1	17542.036083	9135	Buick	Century	Sedan 4D	Sedan	6	3.1	
2	16218.847862	13196	Buick	Century	Sedan 4D	Sedan	6	3.1	
3	16336.913140	16342	Buick	Century	Sedan 4D	Sedan	6	3.1	
4	16339.170324	19832	Buick	Century	Sedan 4D	Sedan	6	3.1	

	Doors	Cruise	Sound	Leather	Model_ord
0	4	1	1	1	10
1	4	1	1	0	10
2	4	1	1	0	10
3	4	1	0	0	10
4	4	1	0	1	10


```
[44]: X = df[['Mileage', 'Model_ord', 'Doors']]
      y = df['Price']
      X1= sm.add_constant(X)
      est = sm.OLS(y,X1).fit()
      est.summary()
```

```
[44]: <class 'statsmodels.iolib.summary.Summary'>
      """
                                OLS Regression Results
=====
Dep. Variable:                  Price    R-squared:                0.042
Model:                            OLS    Adj. R-squared:           0.038
Method:                 Least Squares    F-statistic:                11.57
Date:                  Fri, 21 Jul 2023    Prob (F-statistic):        1.98e-07
Time:                  16:32:46    Log-Likelihood:           -8519.1
No. Observations:                804    AIC:                      1.705e+04
Df Residuals:                    800    BIC:                      1.706e+04
Df Model:                        3
Covariance Type:                nonrobust
=====
                                coef    std err          t      P>|t|      [0.025    0.975]
-----
const                3.125e+04    1809.549     17.272     0.000     2.77e+04     3.48e+04
Mileage               -0.1765         0.042     -4.227     0.000        -0.259     -0.095
Model_ord            -39.0387        39.326     -0.993     0.321     -116.234      38.157
Doors               -1652.9303       402.649     -4.105     0.000    -2443.303    -862.558
=====
Omnibus:                 206.410    Durbin-Watson:           0.080
Prob(Omnibus):            0.000    Jarque-Bera (JB):        470.872
Skew:                     1.379    Prob(JB):                5.64e-103
Kurtosis:                 5.541    Cond. No.                1.15e+05
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
[2] The condition number is large, 1.15e+05. This might indicate that there are
strong multicollinearity or other numerical problems.
      """
```

```
[48]: y.groupby(df.Doors).mean().round(2)
```

```
[48]: Doors
2      23807.14
4      20580.67
Name: Price, dtype: float64
```

```
[86]: import os
import io
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from pandas import DataFrame
```

```
[122]: import warnings
warnings.filterwarnings("ignore")
```

```
[123]: def readFiles(path):
    for root, dirnames, filenames in os.walk(path):
        for filename in filenames:
            path = os.path.join(root, filename)

            inBody = False
            lines = []
            f = io.open(path, 'r', encoding = 'latin1')
            for line in f:
                if inBody:
                    lines.append(line)
                elif line == '\n':
                    inBody = True
            f.close()
            message = '\n'.join(lines)
            yield path, message

def dataframeFromDirectory(path, classification):
    rows = []
    index = []
    for filename, message in readFiles(path):
        rows.append({"message":message, "class":classification})
        index.append(filename)
    return DataFrame(rows, index= index)
```

```
[167]: data = DataFrame({'message':[], 'class': []})
data = data.append(dataFrameFromDirectory('../Rstudio and Python/DataScience/
↳DataScience-Python3/emails/spam', 'spam'))
data = data.append(dataFrameFromDirectory('../Rstudio and Python/DataScience/
↳DataScience-Python3/emails/ham', 'ham'))
```

```
[168]: data.head()
```

```
[168]:          message \
../Rstudio and Python/DataScience/DataScience-P... <!DOCTYPE HTML PUBLIC
"-//W3C//DTD HTML 4.0 Tr...
../Rstudio and Python/DataScience/DataScience-P... 1) Fight The Risk of
Cancer!\n\nhttp://www.adc...
```

```

../Rstudio and Python/DataScience/DataScience-P... 1) Fight The Risk of
Cancer!\n\nhttp://www.adc...
../Rstudio and Python/DataScience/DataScience-P...
#####...
../Rstudio and Python/DataScience/DataScience-P... I thought you might like
these:\n\n1) Slim Dow...

```

```

class
../Rstudio and Python/DataScience/DataScience-P... spam
../Rstudio and Python/DataScience/DataScience-P... spam
../Rstudio and Python/DataScience/DataScience-P... spam
../Rstudio and Python/DataScience/DataScience-P... spam
../Rstudio and Python/DataScience/DataScience-P... spam

```

```

[169]: vectorizer = CountVectorizer()
counts = vectorizer.fit_transform(data['message'].values)
classifier = MultinomialNB()
targets = data['class'].values
classifier.fit(counts, targets)
classifier

```

```

[169]: MultinomialNB()

```

```

[170]: example = ['Free Money now!!', 'Hello']
example_counts = vectorizer.transform(example)
predictions = classifier.predict(example_counts)
predictions

```

```

[170]: array(['spam', 'ham'], dtype='<U4')

```

```

[ ]:

```

```

[ ]:

```